

Johns Hopkins COVID 19 Report

DTSA 5301 Data Science as a Field

MS Data Science, University of Colorado Boulder

2024-08-20

Setup Knit Options

echo = true will display code chunks in the output

Load Libraries

```
library(tidyverse)

## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr     1.1.4      v readr     2.1.5
## vforcats   1.0.0      v stringr   1.5.1
## v ggplot2   3.5.0      v tibble    3.2.1
## v lubridate 1.9.3      v tidyrr    1.3.1
## v purrr    1.0.2

## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()   masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors

library(conflicted)
library(lubridate)
library(caret)

## Loading required package: lattice

library(xgboost)
library(pROC)

## Type 'citation("pROC")' for a citation.

library(PRROC)
library(MLmetrics)
library(glmnet)
```

```
## Loading required package: Matrix
##
## Attaching package: 'Matrix'
##
## The following objects are masked from 'package:tidyverse':
##   expand, pack, unpack
##
## Loaded glmnet 4.1-8
```

```
library(car)
```

```
## Loading required package: carData
```

```
library(smotefamily)
library(ROSE)
```

```
## Loaded ROSE 0.0-4
```

Read Datasets

Download COVID 19 data from Johns Hopkins University, for US and Global cases and deaths.

```
url <- "https://raw.githubusercontent.com/CSSEGISandData/COVID-19/master/csse_covid_19_data/csse_covid_19_time_series.csv"
us_cases <- read_csv(url)
```

```
## Rows: 3342 Columns: 1154
## -- Column specification -----
## Delimiter: ","
## chr    (6): iso2, iso3, Admin2, Province_State, Country_Region, Combined_Key
## dbl (1148): UID, code3, FIPS, Lat, Long_, 1/22/20, 1/23/20, 1/24/20, 1/25/20...
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
url <- "https://raw.githubusercontent.com/CSSEGISandData/COVID-19/master/csse_covid_19_data/csse_covid_19_global.csv"
global_cases <- read_csv(url)
```

```
## Rows: 289 Columns: 1147
## -- Column specification -----
## Delimiter: ","
## chr    (2): Province/State, Country/Region
## dbl (1145): Lat, Long, 1/22/20, 1/23/20, 1/24/20, 1/25/20, 1/26/20, 1/27/20, ...
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```

url <- "https://raw.githubusercontent.com/CSSEGISandData/COVID-19/master/csse_covid_19_data/csse_covid_19_us_deaths.csv"
us_deaths <- read_csv(url)

## Rows: 3342 Columns: 1155
## -- Column specification -----
## Delimiter: ","
## chr    (6): iso2, iso3, Admin2, Province_State, Country_Region, Combined_Key
## dbl (1149): UID, code3, FIPS, Lat, Long_, Population, 1/22/20, 1/23/20, 1/24...
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.

url <- "https://raw.githubusercontent.com/CSSEGISandData/COVID-19/master/csse_covid_19_data/csse_covid_19_global_deaths.csv"
global_deaths <- read_csv(url)

## Rows: 289 Columns: 1147
## -- Column specification -----
## Delimiter: ","
## chr    (2): Province/State, Country/Region
## dbl (1145): Lat, Long, 1/22/20, 1/23/20, 1/24/20, 1/25/20, 1/26/20, 1/27/20, ...
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.

```

Inspect Data

Here we can see the first 5 rows of the data and can page to the right to see all the columns. We start with inspecting us cases.

```
head(us_cases)
```

```

## # A tibble: 6 x 1,154
##       UID iso2 iso3 code3 FIPS Admin2 Province_State Country_Region   Lat
##   <dbl> <chr> <chr> <dbl> <dbl> <chr>     <chr>           <dbl>
## 1 84001001 US    USA     840  1001 Autauga Alabama          US      32.5
## 2 84001003 US    USA     840  1003 Baldwin Alabama         US      30.7
## 3 84001005 US    USA     840  1005 Barbour Alabama        US      31.9
## 4 84001007 US    USA     840  1007 Bibb   Alabama         US      33.0
## 5 84001009 US    USA     840  1009 Blount Alabama        US      34.0
## 6 84001011 US    USA     840  1011 Bullock Alabama       US      32.1
## # i 1,145 more variables: Long_ <dbl>, Combined_Key <chr>, '1/22/20' <dbl>,
## #   '1/23/20' <dbl>, '1/24/20' <dbl>, '1/25/20' <dbl>, '1/26/20' <dbl>,
## #   '1/27/20' <dbl>, '1/28/20' <dbl>, '1/29/20' <dbl>, '1/30/20' <dbl>,
## #   '1/31/20' <dbl>, '2/1/20' <dbl>, '2/2/20' <dbl>, '2/3/20' <dbl>,
## #   '2/4/20' <dbl>, '2/5/20' <dbl>, '2/6/20' <dbl>, '2/7/20' <dbl>,
## #   '2/8/20' <dbl>, '2/9/20' <dbl>, '2/10/20' <dbl>, '2/11/20' <dbl>,
## #   '2/12/20' <dbl>, '2/13/20' <dbl>, '2/14/20' <dbl>, '2/15/20' <dbl>, ...

```

Next we inspect us deaths.

```

head(us_deaths)

## # A tibble: 6 x 1,155
##   UID iso2 iso3 code3 FIPS Admin2 Province_State Country_Region   Lat
##   <dbl> <chr> <chr> <dbl> <dbl> <chr> <chr> <chr> <dbl>
## 1 84001001 US    USA     840  1001 Autauga Alabama      US     32.5
## 2 84001003 US    USA     840  1003 Baldwin Alabama     US     30.7
## 3 84001005 US    USA     840  1005 Barbour Alabama     US     31.9
## 4 84001007 US    USA     840  1007 Bibb   Alabama     US     33.0
## 5 84001009 US    USA     840  1009 Blount Alabama     US     34.0
## 6 84001011 US    USA     840  1011 Bullock Alabama    US     32.1
## # i 1,146 more variables: Long_ <dbl>, Combined_Key <chr>, Population <dbl>,
## #   '1/22/20' <dbl>, '1/23/20' <dbl>, '1/24/20' <dbl>, '1/25/20' <dbl>,
## #   '1/26/20' <dbl>, '1/27/20' <dbl>, '1/28/20' <dbl>, '1/29/20' <dbl>,
## #   '1/30/20' <dbl>, '1/31/20' <dbl>, '2/1/20' <dbl>, '2/2/20' <dbl>,
## #   '2/3/20' <dbl>, '2/4/20' <dbl>, '2/5/20' <dbl>, '2/6/20' <dbl>,
## #   '2/7/20' <dbl>, '2/8/20' <dbl>, '2/9/20' <dbl>, '2/10/20' <dbl>,
## #   '2/11/20' <dbl>, '2/12/20' <dbl>, '2/13/20' <dbl>, '2/14/20' <dbl>, ...

```

Next we inspect global cases.

```

head(global_cases)

## # A tibble: 6 x 1,147
##   'Province/State' 'Country/Region'   Lat  Long '1/22/20' '1/23/20' '1/24/20'
##   <chr>           <chr>       <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 <NA>            Afghanistan  33.9  67.7    0     0     0
## 2 <NA>            Albania     41.2  20.2    0     0     0
## 3 <NA>            Algeria     28.0  1.66    0     0     0
## 4 <NA>            Andorra     42.5  1.52    0     0     0
## 5 <NA>            Angola      -11.2 17.9    0     0     0
## 6 <NA>            Antarctica -71.9  23.3    0     0     0
## # i 1,140 more variables: '1/25/20' <dbl>, '1/26/20' <dbl>, '1/27/20' <dbl>,
## #   '1/28/20' <dbl>, '1/29/20' <dbl>, '1/30/20' <dbl>, '1/31/20' <dbl>,
## #   '2/1/20' <dbl>, '2/2/20' <dbl>, '2/3/20' <dbl>, '2/4/20' <dbl>,
## #   '2/5/20' <dbl>, '2/6/20' <dbl>, '2/7/20' <dbl>, '2/8/20' <dbl>,
## #   '2/9/20' <dbl>, '2/10/20' <dbl>, '2/11/20' <dbl>, '2/12/20' <dbl>,
## #   '2/13/20' <dbl>, '2/14/20' <dbl>, '2/15/20' <dbl>, '2/16/20' <dbl>,
## #   '2/17/20' <dbl>, '2/18/20' <dbl>, '2/19/20' <dbl>, '2/20/20' <dbl>, ...

```

Next we inspect global deaths.

```

head(global_deaths)

## # A tibble: 6 x 1,147
##   'Province/State' 'Country/Region'   Lat  Long '1/22/20' '1/23/20' '1/24/20'
##   <chr>           <chr>       <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 <NA>            Afghanistan  33.9  67.7    0     0     0
## 2 <NA>            Albania     41.2  20.2    0     0     0
## 3 <NA>            Algeria     28.0  1.66    0     0     0
## 4 <NA>            Andorra     42.5  1.52    0     0     0

```

```

## 5 <NA>          Angola      -11.2 17.9      0      0      0
## 6 <NA>        Antarctica -71.9 23.3      0      0      0
## # i 1,140 more variables: '1/25/20' <dbl>, '1/26/20' <dbl>, '1/27/20' <dbl>,
## # '1/28/20' <dbl>, '1/29/20' <dbl>, '1/30/20' <dbl>, '1/31/20' <dbl>,
## # '2/1/20' <dbl>, '2/2/20' <dbl>, '2/3/20' <dbl>, '2/4/20' <dbl>,
## # '2/5/20' <dbl>, '2/6/20' <dbl>, '2/7/20' <dbl>, '2/8/20' <dbl>,
## # '2/9/20' <dbl>, '2/10/20' <dbl>, '2/11/20' <dbl>, '2/12/20' <dbl>,
## # '2/13/20' <dbl>, '2/14/20' <dbl>, '2/15/20' <dbl>, '2/16/20' <dbl>,
## # '2/17/20' <dbl>, '2/18/20' <dbl>, '2/19/20' <dbl>, '2/20/20' <dbl>, ...

```

pivot data

Next we will pivot the data to make it easier to work with. This will transform the data from wide to long format. We start with us cases.

```

# pivot us cases
us_cases <- us_cases %>%
  pivot_longer(cols = -c(UID, iso2, iso3, code3, FIPS, Admin2, Province_State, Country_Region, Lat, Long))
  select(-c(UID, iso2, iso3, code3, FIPS, Lat, Long_, Combined_Key))
us_cases

## # A tibble: 3,819,906 x 5
##   Admin2 Province_State Country_Region date   cases
##   <chr>   <chr>       <chr>        <chr> <dbl>
## 1 Autauga Alabama    US           1/22/20    0
## 2 Autauga Alabama    US           1/23/20    0
## 3 Autauga Alabama    US           1/24/20    0
## 4 Autauga Alabama    US           1/25/20    0
## 5 Autauga Alabama    US           1/26/20    0
## 6 Autauga Alabama    US           1/27/20    0
## 7 Autauga Alabama    US           1/28/20    0
## 8 Autauga Alabama    US           1/29/20    0
## 9 Autauga Alabama    US           1/30/20    0
## 10 Autauga Alabama   US           1/31/20    0
## # i 3,819,896 more rows

```

We do the same for us deaths.

```

# pivot us deaths
us_deaths <- us_deaths %>%
  pivot_longer(cols = -c(UID, iso2, iso3, code3, FIPS, Admin2, Province_State, Country_Region, Lat, Long))
  select(-c(UID, iso2, iso3, code3, FIPS, Lat, Long_))
us_deaths

## # A tibble: 3,819,906 x 7
##   Admin2 Province_State Country_Region Combined_Key   Population date   deaths
##   <chr>   <chr>       <chr>        <chr>        <dbl> <chr> <dbl>
## 1 Autauga Alabama    US           Autauga, Alaba~ 55869 1/22~    0
## 2 Autauga Alabama    US           Autauga, Alaba~ 55869 1/23~    0
## 3 Autauga Alabama    US           Autauga, Alaba~ 55869 1/24~    0
## 4 Autauga Alabama    US           Autauga, Alaba~ 55869 1/25~    0

```

```

## 5 Autauga Alabama      US          Autauga, Alaba~    55869 1/26~    0
## 6 Autauga Alabama      US          Autauga, Alaba~    55869 1/27~    0
## 7 Autauga Alabama      US          Autauga, Alaba~    55869 1/28~    0
## 8 Autauga Alabama      US          Autauga, Alaba~    55869 1/29~    0
## 9 Autauga Alabama      US          Autauga, Alaba~    55869 1/30~    0
## 10 Autauga Alabama     US          Autauga, Alaba~   55869 1/31~    0
## # i 3,819,896 more rows

```

We do the same for global cases.

```

# pivot global cases
global_cases <- global_cases %>%
  pivot_longer(cols = -c("Province/State", "Country/Region", "Lat", "Long"), names_to = "date", values_to = "cases")
  select(-c(Lat, Long))
global_cases

## # A tibble: 330,327 x 4
##   `Province/State` `Country/Region` date     cases
##   <chr>           <chr>        <chr>    <dbl>
## 1 <NA>             Afghanistan  1/22/20    0
## 2 <NA>             Afghanistan  1/23/20    0
## 3 <NA>             Afghanistan  1/24/20    0
## 4 <NA>             Afghanistan  1/25/20    0
## 5 <NA>             Afghanistan  1/26/20    0
## 6 <NA>             Afghanistan  1/27/20    0
## 7 <NA>             Afghanistan  1/28/20    0
## 8 <NA>             Afghanistan  1/29/20    0
## 9 <NA>             Afghanistan  1/30/20    0
## 10 <NA>            Afghanistan  1/31/20    0
## # i 330,317 more rows

```

We do the same for global deaths.

```

# pivot global deaths
global_deaths <- global_deaths %>%
  pivot_longer(cols = -c("Province/State", "Country/Region", "Lat", "Long"), names_to = "date", values_to = "deaths")
  select(-c(Lat, Long))
global_deaths

## # A tibble: 330,327 x 4
##   `Province/State` `Country/Region` date     deaths
##   <chr>           <chr>        <chr>    <dbl>
## 1 <NA>             Afghanistan  1/22/20    0
## 2 <NA>             Afghanistan  1/23/20    0
## 3 <NA>             Afghanistan  1/24/20    0
## 4 <NA>             Afghanistan  1/25/20    0
## 5 <NA>             Afghanistan  1/26/20    0
## 6 <NA>             Afghanistan  1/27/20    0
## 7 <NA>             Afghanistan  1/28/20    0
## 8 <NA>             Afghanistan  1/29/20    0
## 9 <NA>             Afghanistan  1/30/20    0
## 10 <NA>            Afghanistan  1/31/20    0
## # i 330,317 more rows

```

combine cases and deaths, rename columns, convert dates

Next we will combine cases and deaths, rename columns, and convert dates. We can combine the us cases and deaths because they have columns in common to match on. Then we rename the column “Admin2” to “City” to be more descriptive and match the global data. Then we convert the date column to a date object. Finally we select the columns we want to keep. We will start with the us data.

```
us <- us_cases %>%
  full_join(us_deaths) %>%
  rename("City" = "Admin2") %>%
  mutate(date = mdy(date)) %>%
  select("City", "Province_State", "Country_Region", "date", "cases", "deaths", "Population", "Combined_
```



```
## Joining with 'by = join_by(Admin2, Province_State, Country_Region, date)'
```



```
us
```



```
## # A tibble: 3,819,906 x 8
##   City    Province_State Country_Region date      cases  deaths Population
##   <chr>   <chr>        <chr>       <date>     <dbl>   <dbl>     <dbl>
## 1 Autauga Alabama      US          2020-01-22     0       0     55869
## 2 Autauga Alabama      US          2020-01-23     0       0     55869
## 3 Autauga Alabama      US          2020-01-24     0       0     55869
## 4 Autauga Alabama      US          2020-01-25     0       0     55869
## 5 Autauga Alabama      US          2020-01-26     0       0     55869
## 6 Autauga Alabama      US          2020-01-27     0       0     55869
## 7 Autauga Alabama      US          2020-01-28     0       0     55869
## 8 Autauga Alabama      US          2020-01-29     0       0     55869
## 9 Autauga Alabama      US          2020-01-30     0       0     55869
## 10 Autauga Alabama     US          2020-01-31    0       0     55869
## # i 3,819,896 more rows
## # i 1 more variable: Combined_Key <chr>
```

Next we do the same for the global data. We first combine the global cases with deaths. Then we rename the columns to match the us data. Then we convert the date column to a date object. Finally we select the columns we want to keep.

```
global <- global_cases %>%
  full_join(global_deaths) %>%
  rename("Province_State" = "Province/State", "Country_Region" = "Country/Region") %>%
  mutate(date = mdy(date))
```



```
## Joining with 'by = join_by('Province/State', 'Country/Region', date)'
```



```
global
```



```
## # A tibble: 330,327 x 5
##   Province_State Country_Region date      cases  deaths
##   <chr>        <chr>       <date>     <dbl>   <dbl>
## 1 <NA>         Afghanistan 2020-01-22     0       0
## 2 <NA>         Afghanistan 2020-01-23     0       0
```

```

## 3 <NA> Afghanistan 2020-01-24 0 0
## 4 <NA> Afghanistan 2020-01-25 0 0
## 5 <NA> Afghanistan 2020-01-26 0 0
## 6 <NA> Afghanistan 2020-01-27 0 0
## 7 <NA> Afghanistan 2020-01-28 0 0
## 8 <NA> Afghanistan 2020-01-29 0 0
## 9 <NA> Afghanistan 2020-01-30 0 0
## 10 <NA> Afghanistan 2020-01-31 0 0
## # i 330,317 more rows

```

Summary Data

Next we will summarize the data to see the number of cases and deaths. If the features are numeric we get the mean, median, min, max, and quartiles. If the features are categorical we get the count of each category. We start with the us data.

```
summary(us)
```

```

##      City          Province_State    Country_Region        date
##  Length:3819906  Length:3819906  Length:3819906  Min.   :2020-01-22
##  Class :character Class :character  Class :character  1st Qu.:2020-11-02
##  Mode  :character Mode  :character  Mode  :character  Median  :2021-08-15
##                                         Mean   :2021-08-15
##                                         3rd Qu.:2022-05-28
##                                         Max.   :2023-03-09
##      cases         deaths       Population Combined_Key
##  Min.   :-3073   Min.   :-82.0   Min.   :     0  Length:3819906
##  1st Qu.:  330   1st Qu.:  4.0   1st Qu.: 9917  Class :character
##  Median  : 2272   Median : 37.0   Median : 24892  Mode  :character
##  Mean   : 14088   Mean   : 186.9   Mean   : 99604
##  3rd Qu.:  8159   3rd Qu.: 122.0   3rd Qu.: 64979
##  Max.   :3710586   Max.   :35545.0   Max.   :10039107

```

Next we summarize the global data.

```
summary(global)
```

```

##      Province_State    Country_Region        date        cases
##  Length:330327  Length:330327  Min.   :2020-01-22  Min.   :     0
##  Class :character Class :character  1st Qu.:2020-11-02  1st Qu.: 680
##  Mode  :character Mode  :character  Median :2021-08-15  Median : 14429
##                                         Mean   :2021-08-15  Mean   : 959384
##                                         3rd Qu.:2022-05-28  3rd Qu.: 228517
##                                         Max.   :2023-03-09  Max.   :103802702
##      deaths
##  Min.   :     0
##  1st Qu.:     3
##  Median :   150
##  Mean   : 13380
##  3rd Qu.: 3032
##  Max.   :1123836

```

Filter Data

Next we will filter the data to remove rows where cases are 0. We start with the us data.

```
us <- us %>% dplyr::filter(cases > 0)
summary(us)
```

```
##      City      Province_State      Country_Region      date
##  Length:3474292  Length:3474292  Length:3474292  Min.   :2020-01-22
##  Class :character  Class :character  Class :character  1st Qu.:2020-12-27
##  Mode  :character  Mode  :character  Mode  :character  Median  :2021-09-20
##                                         Mean   :2021-09-19
##                                         3rd Qu.:2022-06-15
##                                         Max.   :2023-03-09
##      cases      deaths      Population      Combined_Key
##  Min.   :     1   Min.   :    0.0   Min.   :      0   Length:3474292
##  1st Qu.:   687   1st Qu.:   10.0   1st Qu.:  10953   Class :character
##  Median :  2849   Median :   47.0   Median :  26248   Mode  :character
##  Mean   : 15489   Mean   :  205.1   Mean   : 104502
##  3rd Qu.:  9345   3rd Qu.:  137.0   3rd Qu.: 68098
##  Max.   :3710586   Max.   :35545.0   Max.   :10039107
```

Next we filter the global data.

```
global <- global %>% dplyr::filter(cases > 0)
summary(global)
```

```
##      Province_State      Country_Region      date      cases
##  Length:306827  Length:306827  Min.   :2020-01-22  Min.   :     1
##  Class :character  Class :character  1st Qu.:2020-12-12  1st Qu.: 1316
##  Mode  :character  Mode  :character  Median :2021-09-16  Median : 20365
##                                         Mean   :2021-09-11  Mean   : 1032863
##                                         3rd Qu.:2022-06-15  3rd Qu.: 271281
##                                         Max.   :2023-03-09  Max.   :103802702
##      deaths
##  Min.   :     0
##  1st Qu.:     7
##  Median :   214
##  Mean   : 14405
##  3rd Qu.:  3665
##  Max.   :1123836
```

Create Combined Key

Next we will create a combined key to uniquely identify each row. We will combine the “Province_State” and “Country_Region” columns to match the us data.

```
global <- global %>%
  unite("Combined_Key", c("Province_State", "Country_Region"), sep = " ", na.rm = TRUE, remove = FALSE)
global
```

```

## # A tibble: 306,827 x 6
##   Combined_Key Province_State Country_Region date       cases deaths
##   <chr>        <chr>        <chr>        <date>      <dbl>  <dbl>
## 1 Afghanistan  <NA>         Afghanistan  2020-02-24     5     0
## 2 Afghanistan  <NA>         Afghanistan  2020-02-25     5     0
## 3 Afghanistan  <NA>         Afghanistan  2020-02-26     5     0
## 4 Afghanistan  <NA>         Afghanistan  2020-02-27     5     0
## 5 Afghanistan  <NA>         Afghanistan  2020-02-28     5     0
## 6 Afghanistan  <NA>         Afghanistan  2020-02-29     5     0
## 7 Afghanistan  <NA>         Afghanistan  2020-03-01     5     0
## 8 Afghanistan  <NA>         Afghanistan  2020-03-02     5     0
## 9 Afghanistan  <NA>         Afghanistan  2020-03-03     5     0
## 10 Afghanistan <NA>         Afghanistan  2020-03-04    5     0
## # i 306,817 more rows

```

Fetch Population

Next we will fetch the population data from the UID_ISO_FIPS_LookUp_Table. We will join the population data to the global data only because the us data already has the population data.

```

url <- "https://raw.githubusercontent.com/CSSEGISandData/COVID-19/master/csse_covid_19_data/UID_ISO_FIPS_LookUp_Table.csv"
uid <- read_csv(url) %>%
  select(-c(Lat, Long_, Combined_Key, code3, iso2, iso3, Admin2))

## Rows: 4321 Columns: 12
## -- Column specification -----
## Delimiter: ","
## chr (7): iso2, iso3, FIPS, Admin2, Province_State, Country_Region, Combined_Key
## dbl (5): UID, code3, Lat, Long_, Population
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.

uid

## # A tibble: 4,321 x 5
##   UID FIPS Province_State Country_Region      Population
##   <dbl> <chr> <chr>        <chr>            <dbl>
## 1 4    <NA>  <NA>         Afghanistan      38928341
## 2 8    <NA>  <NA>         Albania          2877800
## 3 10   <NA>  <NA>         Antarctica      NA
## 4 12   <NA>  <NA>         Algeria         43851043
## 5 20   <NA>  <NA>         Andorra          77265
## 6 24   <NA>  <NA>         Angola           32866268
## 7 28   <NA>  <NA>         Antigua and Barbuda 97928
## 8 32   <NA>  <NA>         Argentina       45195777
## 9 51   <NA>  <NA>         Armenia          2963234
## 10 40   <NA>  <NA>         Austria          9006400
## # i 4,311 more rows

```

Next we join the population data to the global data.

```

global <- global %>%
  left_join(uid, by = c("Province_State", "Country_Region")) %>%
  select(-c("UID", "FIPS")) %>%
  select("Province_State", "Country_Region", "date", "cases", "deaths", "Population", "Combined_Key")
global

## # A tibble: 306,827 x 7
##   Province_State Country_Region date      cases  deaths Population Combined_Key
##   <chr>          <chr>       <date>    <dbl>   <dbl>     <dbl> <chr>
## 1 Afghanistan    Afghanistan 2020-02-24     5      0    38928341 Afghanistan
## 2 Afghanistan    Afghanistan 2020-02-25     5      0    38928341 Afghanistan
## 3 Afghanistan    Afghanistan 2020-02-26     5      0    38928341 Afghanistan
## 4 Afghanistan    Afghanistan 2020-02-27     5      0    38928341 Afghanistan
## 5 Afghanistan    Afghanistan 2020-02-28     5      0    38928341 Afghanistan
## 6 Afghanistan    Afghanistan 2020-02-29     5      0    38928341 Afghanistan
## 7 Afghanistan    Afghanistan 2020-03-01     5      0    38928341 Afghanistan
## 8 Afghanistan    Afghanistan 2020-03-02     5      0    38928341 Afghanistan
## 9 Afghanistan    Afghanistan 2020-03-03     5      0    38928341 Afghanistan
## 10 Afghanistan   Afghanistan 2020-03-04    5      0    38928341 Afghanistan
## # i 306,817 more rows

```

visualize data

Next we will visualize the data. We will start with the us data. We will plot the number of cases and deaths over time. We create new columns “cases_per_million” and “deaths_per_million” to normalize the data by population. We are also grouping the us data by Province_State, Country_Region, and date.

```

us_by_state <- us %>%
  group_by(Province_State, Country_Region, date) %>%
  summarize(cases = sum(cases), deaths = sum(deaths), Population = sum(Population)) %>%
  mutate(cases_per_million = cases * 1000000 / Population, deaths_per_million = deaths * 1000000 / Population)
  select(Province_State, Country_Region, date, cases, deaths, cases_per_million, deaths_per_million, Population)
  ungroup()

## `summarise()` has grouped output by 'Province_State', 'Country_Region'. You can
## override using the '.groups' argument.

us_by_state

## # A tibble: 63,216 x 8
##   Province_State Country_Region date      cases  deaths cases_per_million
##   <chr>          <chr>       <date>    <dbl>   <dbl>           <dbl>
## 1 Alabama        US          2020-03-11     3      0            5.61
## 2 Alabama        US          2020-03-12     4      0            4.41
## 3 Alabama        US          2020-03-13     8      0            4.86
## 4 Alabama        US          2020-03-14    15      0            9.10
## 5 Alabama        US          2020-03-15    28      0           12.4
## 6 Alabama        US          2020-03-16    36      0           16.0
## 7 Alabama        US          2020-03-17    51      0           20.8
## 8 Alabama        US          2020-03-18    61      0           24.1

```

```

##  9 Alabama      US        2020-03-19    88     0       28.1
## 10 Alabama      US        2020-03-20   115     0       34.5
## # i 63,206 more rows
## # i 2 more variables: deaths_per_million <dbl>, Population <dbl>

```

Next we do something similar for the us totals. We take the us data by state which includes the new normalized metrics for cases per million and deaths per million, then group by Country_Region and date, and summarize the cases and deaths. Finally we select the columns we want to keep.

```

us_totals <- us_by_state %>%
  group_by(Country_Region, date) %>%
  summarize(cases = sum(cases), deaths = sum(deaths), Population = sum(Population)) %>%
  mutate(cases_per_million = cases * 1000000 / Population, deaths_per_million = deaths * 1000000 / Population) %>%
  select(Country_Region, date, cases, deaths, cases_per_million, deaths_per_million, Population) %>%
  ungroup()

```

```

## `summarise()` has grouped output by 'Country_Region'. You can override using
## the 'groups' argument.

```

```
us_totals
```

```

## # A tibble: 1,143 x 7
##   Country_Region date      cases  deaths cases_per_million deaths_per_million
##   <chr>        <date>    <dbl>   <dbl>        <dbl>            <dbl>
## 1 US          2020-01-22     1      0        0.444             0
## 2 US          2020-01-23     1      0        0.444             0
## 3 US          2020-01-24     2      0        0.270             0
## 4 US          2020-01-25     2      0        0.270             0
## 5 US          2020-01-26     5      0        0.199             0
## 6 US          2020-01-27     5      0        0.199             0
## 7 US          2020-01-28     5      0        0.199             0
## 8 US          2020-01-29     6      0        0.232             0
## 9 US          2020-01-30     6      0        0.232             0
## 10 US         2020-01-31    8      0        0.287             0
## # i 1,133 more rows
## # i 1 more variable: Population <dbl>

```

Next we plot the us data. We plot the number of cases and deaths over time. We use a log scale for the y axis to better visualize the data. We can see that the number of cases and deaths are increasing over time and have a similar trend.

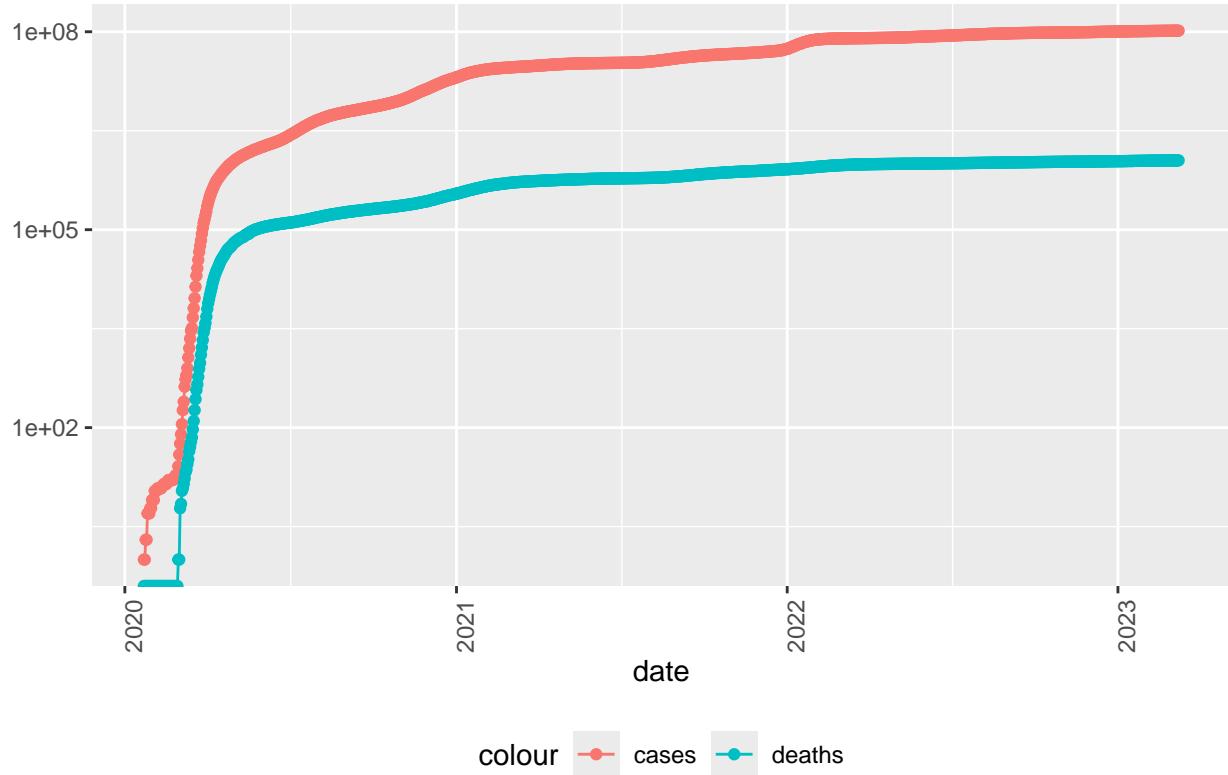
```

us_totals %>%
  dplyr::filter(cases > 0) %>%
  ggplot(aes(x = date, y = cases)) +
  geom_line(aes(color = "cases")) +
  geom_point(aes(color = "cases")) +
  geom_line(aes(y = deaths, color = "deaths")) +
  geom_point(aes(y = deaths, color = "deaths")) +
  scale_y_log10() +
  theme(legend.position = "bottom", axis.text.x = element_text(angle = 90)) +
  labs(title = "COVID19 in US", y = NULL)

```

```
## Warning in scale_y_log10(): log-10 transformation introduced infinite values.
## log-10 transformation introduced infinite values.
```

COVID19 in US

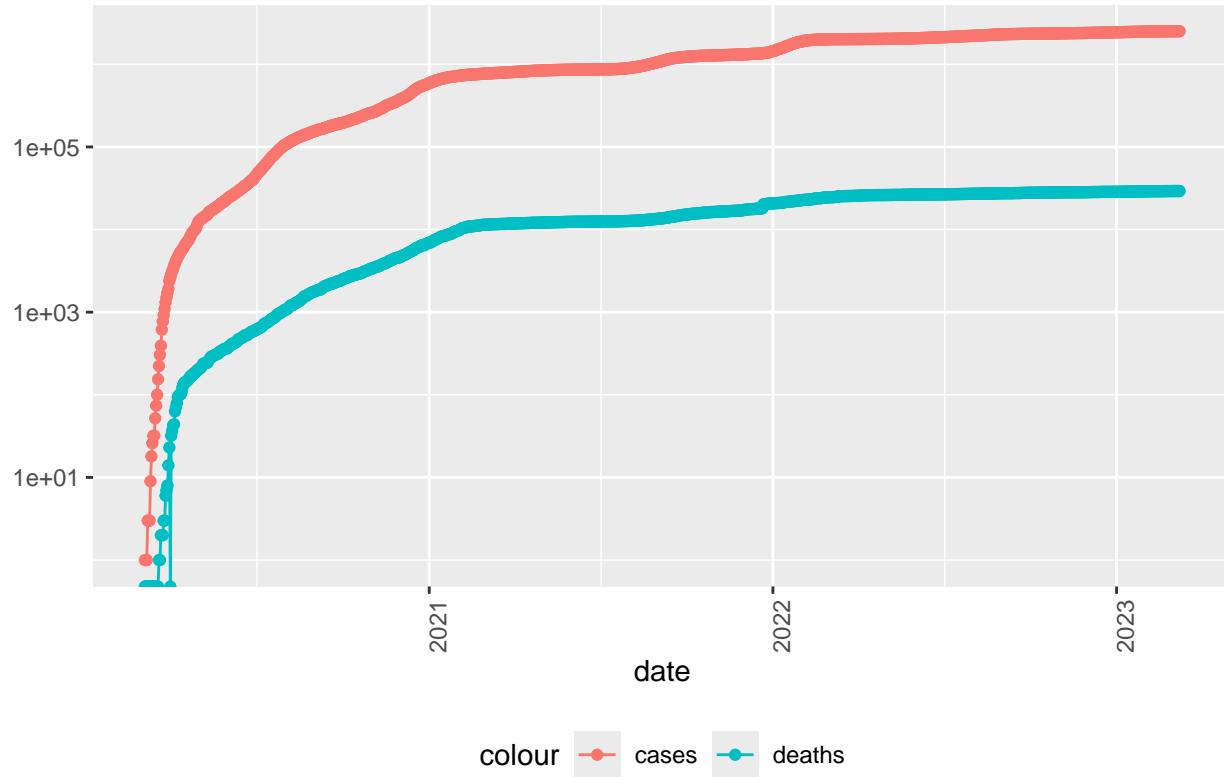


Next we filter down to just one state Tennessee and plot the data. We can see that the number of cases and deaths are increasing over time and have a similar trend. We can also see the trends are similar to the us total data. This would be an indication that the state of Tennessee is following the same trend as the US as a whole and is not an outlier, either higher or lower in rates.

```
state <- "Tennessee"
us_by_state %>%
  dplyr::filter(Province_State == state) %>%
  dplyr::filter(cases > 0) %>%
  ggplot(aes(x = date, y = cases)) +
  geom_line(aes(color = "cases")) +
  geom_point(aes(color = "cases")) +
  geom_line(aes(y = deaths, color = "deaths")) +
  geom_point(aes(y = deaths, color = "deaths")) +
  scale_y_log10() +
  theme(legend.position = "bottom", axis.text.x = element_text(angle = 90)) +
  labs(title = str_c("COVID19 in ", state), y = NULL)
```

```
## Warning in scale_y_log10(): log-10 transformation introduced infinite values.
## log-10 transformation introduced infinite values.
```

COVID19 in Tennessee



Here we check for the maximum date, cases, and deaths in the us data.

```
max(us_totals$date)
```

```
## [1] "2023-03-09"
```

```
max(us_totals$cases)
```

```
## [1] 103802702
```

```
max(us_totals$deaths)
```

```
## [1] 1122724
```

Analyze Data

Next we will analyze the data. We will start by creating new features for new cases and new deaths. This will be the difference between the current and previous day.

```
us <- us %>%
  mutate(new_cases = cases - dplyr::lag(cases), new_deaths = deaths - dplyr::lag(deaths))
us_by_state <- us_by_state %>%
```

```

    mutate(new_cases = cases - dplyr::lag(cases), new_deaths = deaths - dplyr::lag(deaths))
us_totals <- us_totals %>%
    mutate(new_cases = cases - dplyr::lag(cases), new_deaths = deaths - dplyr::lag(deaths))

```

We can validate below that the new cases and new deaths are being calculated correctly for the us state data by city.

us

```

## # A tibble: 3,474,292 x 10
##   City      Province_State Country_Region date       cases  deaths Population
##   <chr>     <chr>        <chr>        <date>     <dbl>   <dbl>     <dbl>
## 1 Autauga  Alabama        US  2020-03-24     1      0     55869
## 2 Autauga  Alabama        US  2020-03-25     5      0     55869
## 3 Autauga  Alabama        US  2020-03-26     6      0     55869
## 4 Autauga  Alabama        US  2020-03-27     6      0     55869
## 5 Autauga  Alabama        US  2020-03-28     6      0     55869
## 6 Autauga  Alabama        US  2020-03-29     6      0     55869
## 7 Autauga  Alabama        US  2020-03-30     8      0     55869
## 8 Autauga  Alabama        US  2020-03-31     8      0     55869
## 9 Autauga  Alabama        US  2020-04-01    10      0     55869
## 10 Autauga Alabama        US  2020-04-02    12      0     55869
## # i 3,474,282 more rows
## # i 3 more variables: Combined_Key <chr>, new_cases <dbl>, new_deaths <dbl>

```

We can validate below that the new cases and new deaths are being calculated correctly for the us state totals.

us_by_state

```

## # A tibble: 63,216 x 10
##   Province_State Country_Region date       cases  deaths cases_per_million
##   <chr>        <chr>        <date>     <dbl>   <dbl>          <dbl>
## 1 Alabama        US  2020-03-11     3      0          5.61
## 2 Alabama        US  2020-03-12     4      0          4.41
## 3 Alabama        US  2020-03-13     8      0          4.86
## 4 Alabama        US  2020-03-14    15      0          9.10
## 5 Alabama        US  2020-03-15    28      0         12.4
## 6 Alabama        US  2020-03-16    36      0         16.0
## 7 Alabama        US  2020-03-17    51      0         20.8
## 8 Alabama        US  2020-03-18    61      0         24.1
## 9 Alabama        US  2020-03-19   88      0         28.1
## 10 Alabama       US  2020-03-20   115      0         34.5
## # i 63,206 more rows
## # i 4 more variables: deaths_per_million <dbl>, Population <dbl>,
## #   new_cases <dbl>, new_deaths <dbl>

```

We can validate below that the new cases and new deaths are being calculated correctly for the us totals.

us_totals

```

## # A tibble: 1,143 x 9
##   Country_Region date      cases  deaths cases_per_million deaths_per_million
##   <chr>        <date>    <dbl>   <dbl>       <dbl>            <dbl>
## 1 US          2020-01-22     1      0       0.444            0
## 2 US          2020-01-23     1      0       0.444            0
## 3 US          2020-01-24     2      0       0.270            0
## 4 US          2020-01-25     2      0       0.270            0
## 5 US          2020-01-26     5      0       0.199            0
## 6 US          2020-01-27     5      0       0.199            0
## 7 US          2020-01-28     5      0       0.199            0
## 8 US          2020-01-29     6      0       0.232            0
## 9 US          2020-01-30     6      0       0.232            0
## 10 US         2020-01-31    8      0       0.287            0
## # i 1,133 more rows
## # i 3 more variables: Population <dbl>, new_cases <dbl>, new_deaths <dbl>

```

Next we will calculate the new cases and new deaths per thousand for the us totals.

```

us_totals %>%
  ggplot(aes(x = date, y = new_cases)) +
  geom_line(aes(color = "new_cases")) +
  geom_point(aes(color = "new_cases")) +
  geom_line(aes(y = new_deaths, color = "new_deaths")) +
  geom_point(aes(y = new_deaths, color = "new_deaths")) +
  scale_y_log10() +
  theme(legend.position = "bottom", axis.text.x = element_text(angle = 90)) +
  labs(title = "COVID19 in US", y = NULL)

## Warning in transformation$transform(x): NaNs produced

## Warning in scale_y_log10(): log-10 transformation introduced infinite values.

## Warning in transformation$transform(x): NaNs produced

## Warning in scale_y_log10(): log-10 transformation introduced infinite values.

## Warning in transformation$transform(x): NaNs produced

## Warning in scale_y_log10(): log-10 transformation introduced infinite values.

## Warning in transformation$transform(x): NaNs produced

## Warning in scale_y_log10(): log-10 transformation introduced infinite values.

## Warning: Removed 1 row containing missing values or values outside the scale range
## ('geom_line()').

## Warning: Removed 2 rows containing missing values or values outside the scale range
## ('geom_point()').

```

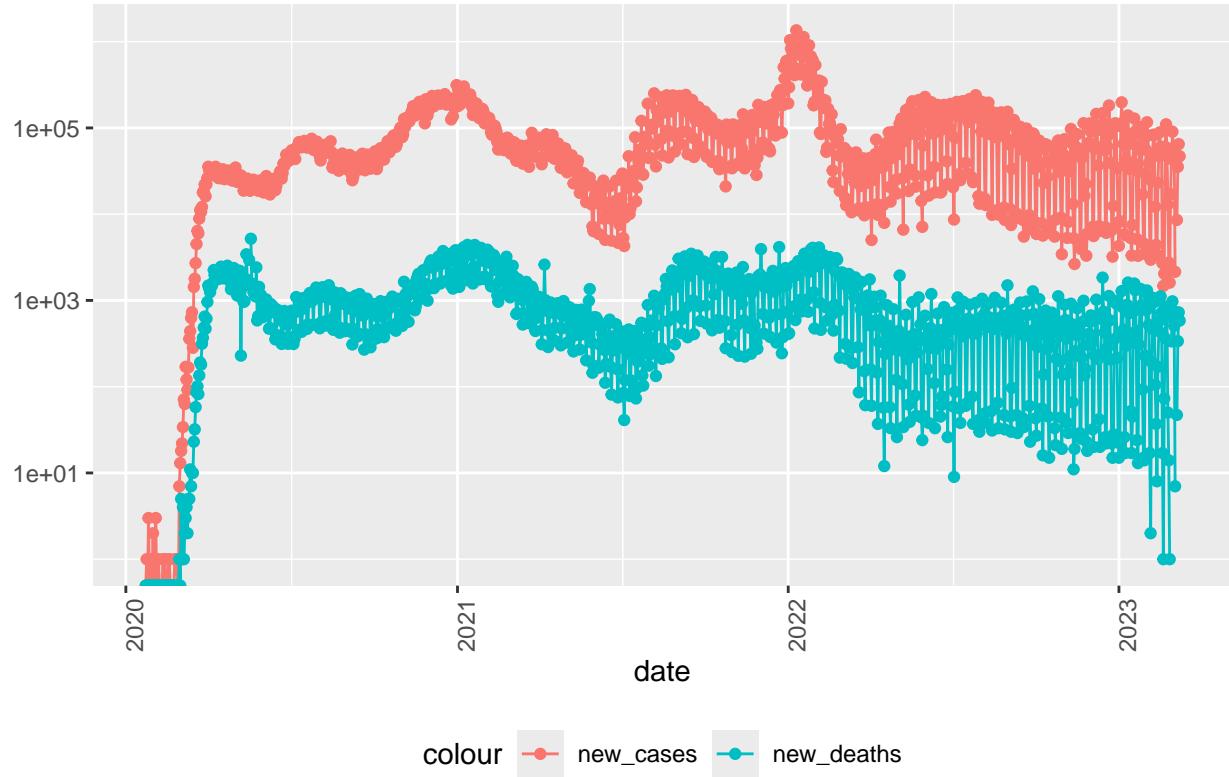
```

## Warning: Removed 1 row containing missing values or values outside the scale range
## ('geom_line()').

## Warning: Removed 7 rows containing missing values or values outside the scale range
## ('geom_point()').

```

COVID19 in US



We do the same for us by state and filter for Tennessee.

```

state <- "Tennessee"
us_by_state %>%
  dplyr::filter(Province_State == state) %>%
  ggplot(aes(x = date, y = new_cases)) +
  geom_line(aes(color = "new_cases")) +
  geom_point(aes(color = "new_cases")) +
  geom_line(aes(y = new_deaths, color = "new_deaths")) +
  geom_point(aes(y = new_deaths, color = "new_deaths")) +
  scale_y_log10() +
  theme(legend.position = "bottom", axis.text.x = element_text(angle = 90)) +
  labs(title = str_c("COVID19 in", state), y = NULL)

```

```

## Warning in transformation$transform(x): NaNs produced

## Warning in scale_y_log10(): log-10 transformation introduced infinite values.

## Warning in transformation$transform(x): NaNs produced

```

```

## Warning in scale_y_log10(): log-10 transformation introduced infinite values.

## Warning in transformation$transform(x): NaNs produced

## Warning in scale_y_log10(): log-10 transformation introduced infinite values.

## Warning in transformation$transform(x): NaNs produced

## Warning in scale_y_log10(): log-10 transformation introduced infinite values.

## Warning: Removed 1 row containing missing values or values outside the scale range
## ('geom_line()').

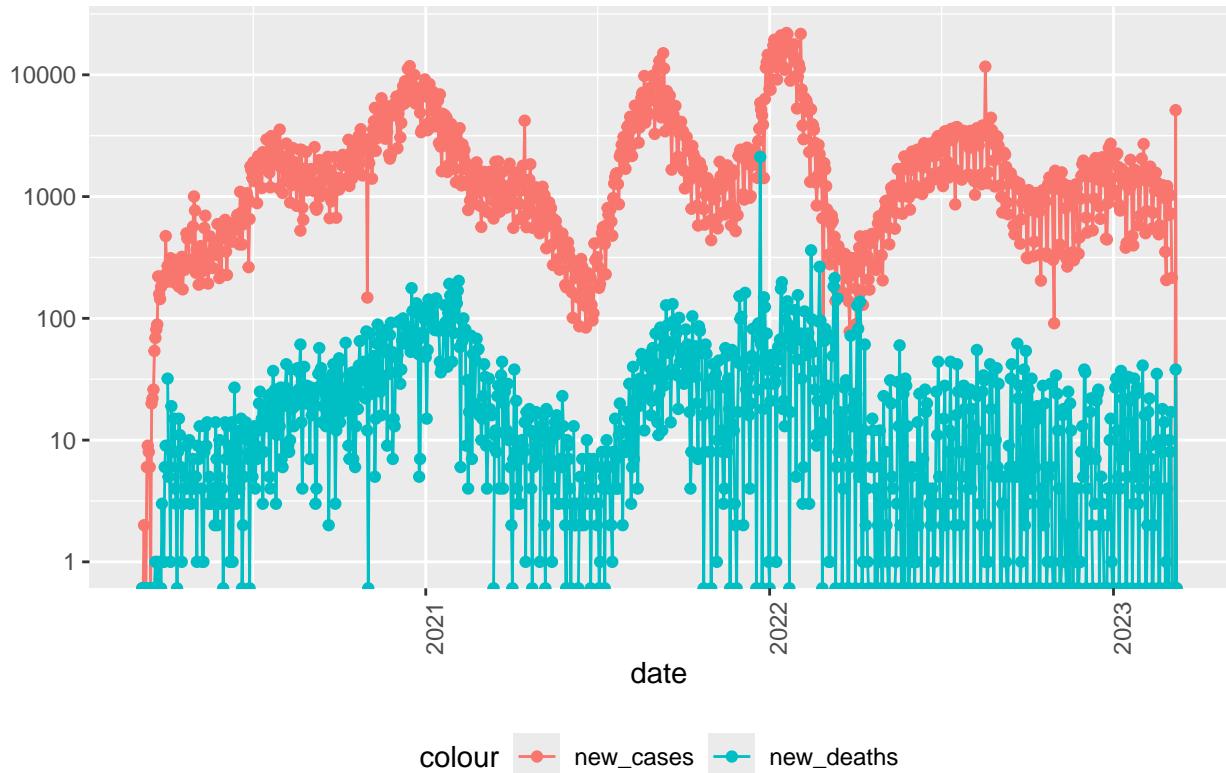
## Warning: Removed 2 rows containing missing values or values outside the scale range
## ('geom_point()').

## Warning: Removed 1 row containing missing values or values outside the scale range
## ('geom_line()').

## Warning: Removed 14 rows containing missing values or values outside the scale range
## ('geom_point()').

```

COVID19 in Tennessee



Next we will summarize the us state totals, and filter for only cases and population greater than 0.

```

us_state_totals <- us_by_state %>%
  group_by(Province_State) %>%
  summarize(cases = max(cases), deaths = max(deaths), Population = max(Population), cases_per_thousand =
  dplyr::filter(cases > 0, Population > 0) %>%
  ungroup()

```

Here we inspect the top 10 min deaths per thousand.

```

us_state_totals %>%
  slice_min(deaths_per_thousand, n = 10)

```

```

## # A tibble: 10 x 6
##   Province_State      cases  deaths Population cases_per_thousand
##   <chr>          <dbl>   <dbl>     <dbl>            <dbl>
## 1 American Samoa      8320     34     55641            150.
## 2 Northern Mariana Islands 13666     41     55144            248.
## 3 Virgin Islands      24813    130    107268            231.
## 4 Hawaii              380608   1841   1415872            269.
## 5 Vermont              152618    929    623989            245.
## 6 Puerto Rico          1101469   5823   3754939            293.
## 7 Utah                 1090346   5298   2785478            391.
## 8 District of Columbia 177945   1432    705749            252.
## 9 Alaska                307655   1486    728809            422.
## 10 Washington           1928913  15683   7614893            253.
## # i 1 more variable: deaths_per_thousand <dbl>

```

Here we inspect the top 10 min deaths per thousand and select the columns we want to keep.

```

us_state_totals %>%
  slice_min(deaths_per_thousand, n = 10) %>%
  select(deaths_per_thousand, cases_per_thousand, everything())

```

```

## # A tibble: 10 x 6
##   deaths_per_thousand cases_per_thousand Province_State      cases  deaths
##   <dbl>          <dbl> <chr>          <dbl>   <dbl>
## 1 0.611            0.32e3 American Samoa      8.32e3   34
## 2 0.744            1.37e4 Northern Mariana Islands 1.37e4   41
## 3 1.21             2.48e4 Virgin Islands      2.48e4  130
## 4 1.30             3.81e5 Hawaii              3.81e5  1841
## 5 1.49             1.53e5 Vermont              1.53e5  929
## 6 1.55             1.10e6 Puerto Rico          1.10e6  5823
## 7 1.90             1.09e6 Utah                 1.09e6  5298
## 8 2.03             1.78e5 District of Columbia 1.78e5  1432
## 9 2.04             3.08e5 Alaska               3.08e5  1486
## 10 2.06            1.93e6 Washington           1.93e6  15683
## # i 1 more variable: Population <dbl>

```

Here we inspect the top 10 max deaths per thousand.

```

us_state_totals %>%
  slice_max(deaths_per_thousand, n = 10)

```

```

## # A tibble: 10 x 6
##   Province_State   cases  deaths Population cases_per_thousand
##   <chr>        <dbl>  <dbl>      <dbl>            <dbl>
## 1 Arizona        2443514 33102    7278717            336.
## 2 Oklahoma       1290929 17972    3956971            326.
## 3 Mississippi    990756  13370    2976149            333.
## 4 West Virginia  642760  7960     1792147            359.
## 5 New Mexico     670929  9061     2096829            320.
## 6 Arkansas        1006883 13020    3017804            334.
## 7 Alabama         1644533 21032    4903185            335.
## 8 Tennessee       2515130 29263    6829174            368.
## 9 Michigan        3064125 42205    9986857            307.
## 10 Kentucky       1718471 18130    4467673            385.
## # i 1 more variable: deaths_per_thousand <dbl>

```

Here we inspect the top 10 max deaths per thousand and select the columns we want to keep.

```

us_state_totals %>%
  slice_max(deaths_per_thousand, n = 10) %>%
  select(deaths_per_thousand, cases_per_thousand, everything())

```

```

## # A tibble: 10 x 6
##   deaths_per_thousand cases_per_thousand Province_State   cases  deaths
##   <dbl>              <dbl>          <chr>        <dbl>  <dbl>
## 1 4.55               336. Arizona     2443514 33102
## 2 4.54               326. Oklahoma    1290929 17972
## 3 4.49               333. Mississippi  990756  13370
## 4 4.44               359. West Virginia 642760  7960
## 5 4.32               320. New Mexico   670929  9061
## 6 4.31               334. Arkansas    1006883 13020
## 7 4.29               335. Alabama     1644533 21032
## 8 4.28               368. Tennessee   2515130 29263
## 9 4.23               307. Michigan    3064125 42205
## 10 4.06              385. Kentucky    1718471 18130
## # i 1 more variable: Population <dbl>

```

Model Data

Next we will model the data. We will start by creating a linear regression model to predict deaths per thousand based on cases per thousand. We will use the us state totals data for this analysis. Lastly we will summarize the model's performance.

```

model <- lm(deaths_per_thousand ~ cases_per_thousand, data = us_state_totals)
summary(model)

```

```

##
## Call:
## lm(formula = deaths_per_thousand ~ cases_per_thousand, data = us_state_totals)
##
## Residuals:
##      Min       1Q   Median       3Q      Max 
## -1.1000  -0.4000   0.0000  -0.1000  1.1000 
## 
```

```

## -2.2394 -0.6114  0.1965  0.6413  1.2413
##
## Coefficients:
##                         Estimate Std. Error t value Pr(>|t|)
## (Intercept)      -0.02599    0.72442  -0.036   0.972
## cases_per_thousand 0.01020    0.00231   4.414 4.89e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.8803 on 54 degrees of freedom
## Multiple R-squared:  0.2652, Adjusted R-squared:  0.2516
## F-statistic: 19.49 on 1 and 54 DF,  p-value: 4.894e-05

```

Next we will predict the deaths per thousand based on the cases per thousand.

```
us_state_totals %>% mutate(pred = predict(model))
```

```

## # A tibble: 56 x 7
##   Province_State     cases  deaths Population cases_per_thousand
##   <chr>        <dbl>   <dbl>      <dbl>            <dbl>
## 1 Alabama       1644533  21032  4903185          335.
## 2 Alaska        307655   1486  728809           422.
## 3 American Samoa     8320    34  55641            150.
## 4 Arizona       2443514  33102  7278717          336.
## 5 Arkansas      1006883  13020  3017804          334.
## 6 California    12129699 101159  39512223          307.
## 7 Colorado       1764401  14181  5758736          306.
## 8 Connecticut    976657   12220  3565287          274.
## 9 Delaware       330793   3324  973764           340.
## 10 District of Columbia 177945   1432  705749          252.
## # i 46 more rows
## # i 2 more variables: deaths_per_thousand <dbl>, pred <dbl>

```

Next we will add the predicted deaths per thousand to the us state totals data.

```
us_totals_w_pred <- us_state_totals %>% mutate(pred = predict(model))
us_totals_w_pred
```

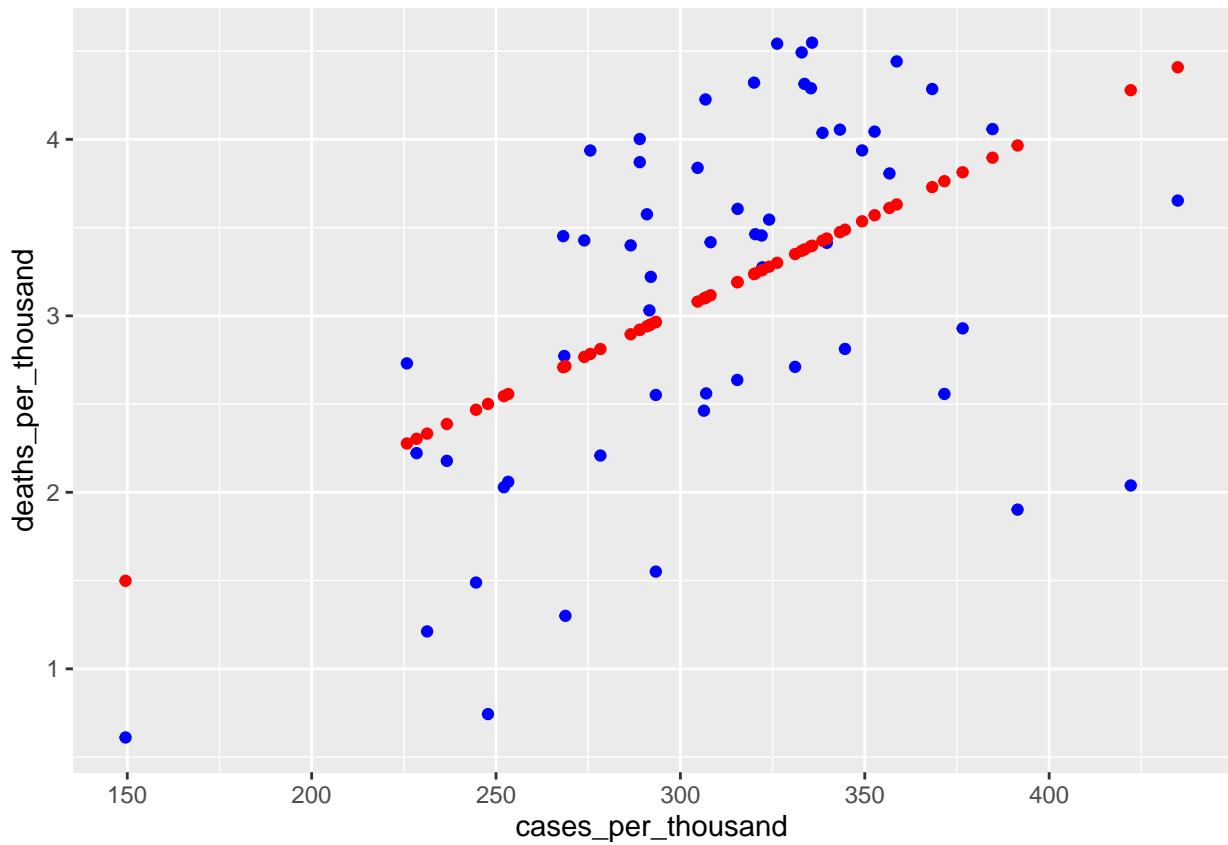
```

## # A tibble: 56 x 7
##   Province_State     cases  deaths Population cases_per_thousand
##   <chr>        <dbl>   <dbl>      <dbl>            <dbl>
## 1 Alabama       1644533  21032  4903185          335.
## 2 Alaska        307655   1486  728809           422.
## 3 American Samoa     8320    34  55641            150.
## 4 Arizona       2443514  33102  7278717          336.
## 5 Arkansas      1006883  13020  3017804          334.
## 6 California    12129699 101159  39512223          307.
## 7 Colorado       1764401  14181  5758736          306.
## 8 Connecticut    976657   12220  3565287          274.
## 9 Delaware       330793   3324  973764           340.
## 10 District of Columbia 177945   1432  705749          252.
## # i 46 more rows
## # i 2 more variables: deaths_per_thousand <dbl>, pred <dbl>

```

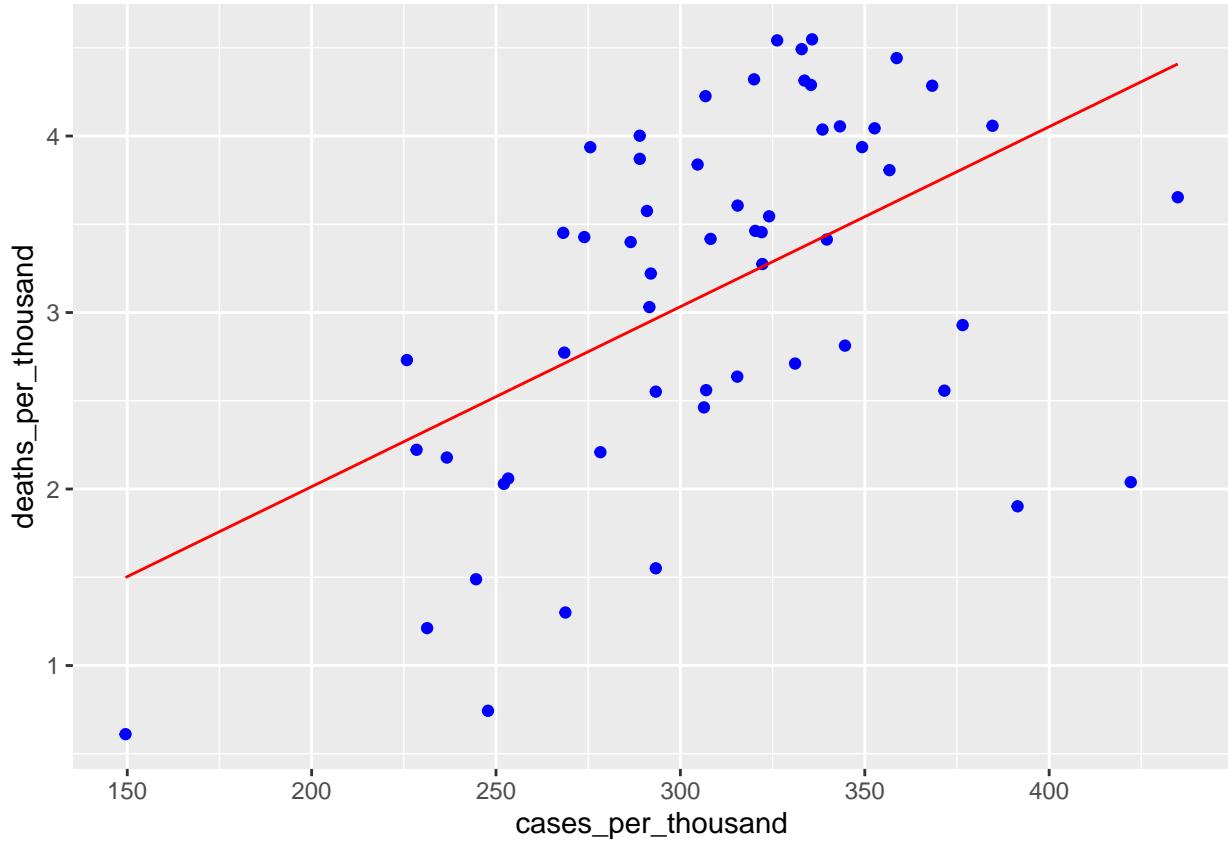
Next we will plot the actual deaths per thousand vs the predicted deaths per thousand. We can see that the predicted deaths per thousand are close to the actual deaths per thousand. The red points are the predicted deaths per thousand and the blue points are the actual deaths per thousand. The blue dots do appear to have a linear relationship with the red dots.

```
us_totals_w_pred %>% ggplot() +
  geom_point(aes(x = cases_per_thousand, y = deaths_per_thousand), color = "blue") +
  geom_point(aes(x = cases_per_thousand, y = pred), color = "red")
```



Next we create another plot to visualize the model.

```
us_totals_w_pred %>% ggplot() +
  geom_point(aes(x = cases_per_thousand, y = deaths_per_thousand), color = "blue") +
  geom_line(aes(x = cases_per_thousand, y = pred), color = "red")
```



Next we summarize the model's performance using RMSE and R2.

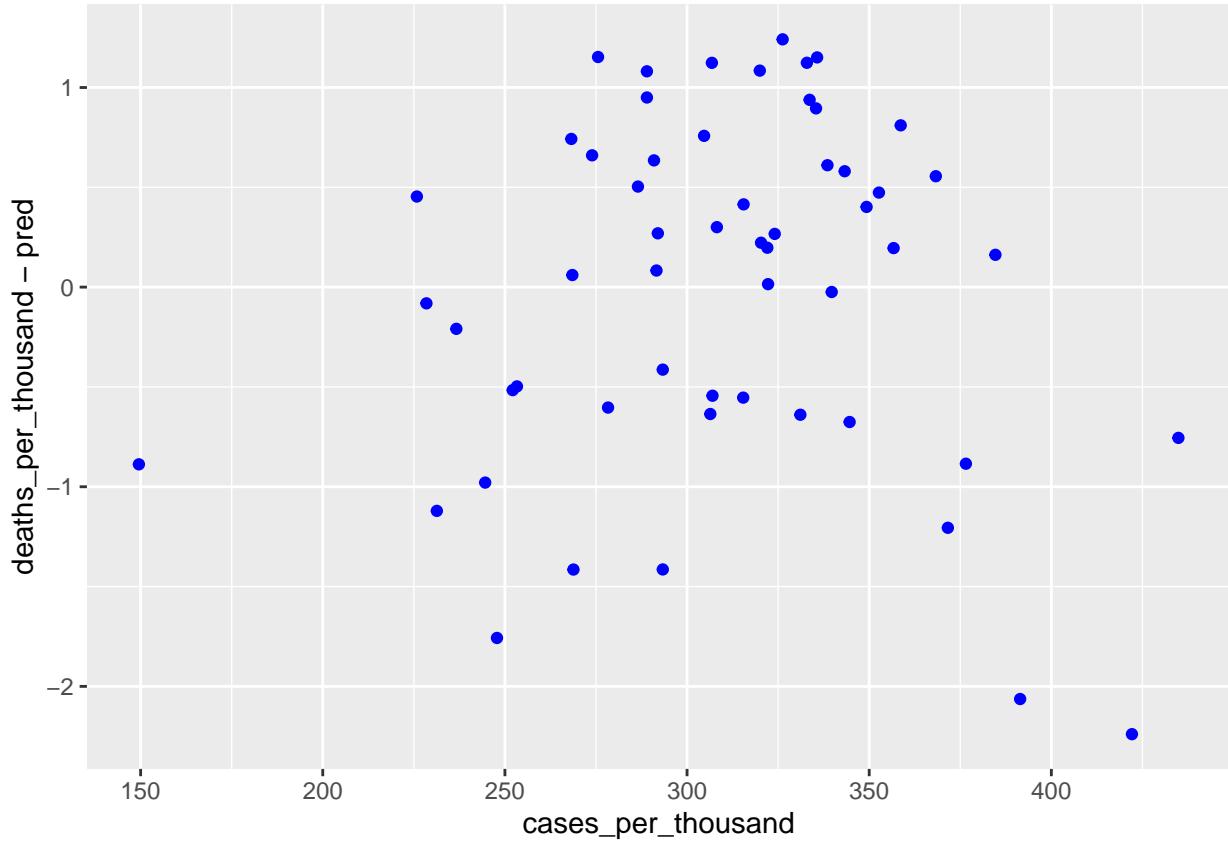
```
us_totals_w_pred %>% summarize(rmse = sqrt(mean((deaths_per_thousand - pred)^2)), r2 = cor(deaths_per_thousand, pred))
```

```
## # A tibble: 1 x 2
##       rmse     r2
##     <dbl> <dbl>
## 1 0.864 0.265
```

Above we can see that the RMSE is 0.8644585 and the R2 is 0.2651695. The RMSE is the square root of the mean of the squared differences between the actual and predicted deaths per thousand. The R2 is the square of the correlation between the actual and predicted deaths per thousand. The RMSE is a measure of the model's accuracy and the R2 is a measure of the model's goodness of fit. An RMSE of 0.8644585 means that the model's predictions are on average 0.8644585 deaths per thousand away from the actual deaths per thousand. An R2 of 0.2651695 means that the model explains 26.52% of the variance in the deaths per thousand.

Here we plot the residuals which are the difference between the actual deaths per thousand and the predicted deaths per thousand. We can see that the residuals are randomly distributed and there is no clear pattern. This is a good indication that the model is a good fit for the data.

```
us_totals_w_pred %>% ggplot() +
  geom_point(aes(x = cases_per_thousand, y = deaths_per_thousand - pred), color = "blue")
```



Conclusion

In conclusion, we have analyzed the COVID 19 data from Johns Hopkins University. We have cleaned the data, visualized the data, created new features, and modeled the data. We then created a linear regression model to predict deaths per thousand based on cases per thousand. We evaluated the model using RMSE which is a measure of the model's accuracy and with R² which is a measure of the model's goodness of fit. An RMSE of 0.8644585 means that the model's predictions are on average 0.8644585 deaths per thousand away from the actual deaths per thousand. An R² of 0.2651695 means that the model explains 26.52% of the variance in the deaths per thousand. The residuals are also randomly distributed and there is no clear pattern, which is a good indication that the model is a good fit for the data. Some bias that could be present is in the data itself, as the data is from Johns Hopkins University and may not be representative of the entire population. We also do not have enough information to determine how positive cases and positive deaths are being reported, or under reported. Therefore, we should be cautious in interpreting the results and making predictions until we have more information. However, this model seems to be a good fit for the data in making accurate predictions for deaths per thousand based on cases per thousand in the US state totals data.